

ECE 459: Programming for Performance


Assignment 1

Patrick Lam*

January 15, 2013 (Due: January 28, 2013)

Background

For this assignment, we'll be parallelizing the **Monte Carlo Estimation for Pi**. It's a crude estimation that works by generating a bunch of random numbers between $[0, 1]$. Consider a 1×1 square with a quarter circle enclosed within.



We can estimate π by picking a collection of random points and counting how many are contained within the quarter circle. We pick two random numbers x and y , check if $x^2 + y^2 \leq 1$, and if it does, then we add 1 to the accumulator. We'll pick i pairs of random points. Let c pairs fall within the quarter circle. Then, we can estimate π by calculating $4 \cdot \frac{c}{i}$ (since the ratio of the quarter circle's area to that of the square is $\frac{\pi}{4}$). Your program will generate i random points, determine c by calculating how many are enclosed in the quarter circle, and finally output π .

Setup

Download the provided assignment code from <http://patricklam.ca/p4p/assignments/provided-assignment-01.tar.gz> and untar it using the command: `tar xzvf assignment-01.tar.gz`.

You should do this assignment in Linux, as the provided Makefile was only tested on Linux and isn't very robust. Since we're testing parallel execution, you may not want to use a virtual machine and instead run the code natively—most free ones don't support multiple CPUs well. You may use the course machine, `ece459-1.uwaterloo.ca`. E-mail Xavier (xnombis@engmail.uwaterloo.ca) from your `uwaterloo.ca` address for `ssh` access. We'll make sure that everything will work on this system. Just in case you want to make sure your setup is the same, I used `gcc 4.7.2` and `glibc 2.13`. Make sure to use the optimization flag `-O2`, or you may get some weird results.

*Assignment thanks to Jon Eyolfson.

Part 1 - Programming (60 marks)

Use the `pthread` library to create a threaded version of the provided program. All of your parallel code should be inside `#ifdef PARALLEL` sections. Your program should create as many threads as specified by `num.threads` (whose value comes from the `-t` command line option). As a reminder, you should make sure all of your external calls are *thread-safe*. For example, `man 3 rand` shows the documentation for `rand`, allowing you to find the thread-safe version. **Note: `rand` is a lot slower than `rand_r`, so the sequential version already uses `rand_r` to ensure that your time comparisons make sense. You should still get used to looking up *thread-safe* calls.** Also, do not change the fundamental algorithm (as bad as it may be), since we want the same amount of work done in the sequential and parallel versions.

Part 2 - Benchmarking (40 marks)

Run the sequential version with the number of iterations (as set with the `-i` option) sufficient for the program to run for at least 10 seconds. You may use the `time` command to benchmark your program. Run it at least 6 times, record the times, and compute the average. Assuming the code is 100% parallel, predict its runtime running on the number of physical cores in your system (4 for `ece459-1`; should be at least 2 in your solution). Explicitly state the number of physical cores in the machine. Next, run your parallel version with the same number of iterations and the number of threads set to the number of physical cores in your system. Run this at least 6 times, record results, and compute the average. Is your parallel runtime approximately equal to your predicted runtime? Write a few sentences describing why or why not your results agree with what you predicted.

Finally, if your CPU has hyperthreading, set the number of threads equal to the number of **virtual CPUs**, run your benchmark 6 times again, record times, and compute the average. Calculate the speedup and verify that it is less than the number of virtual CPUs. Now, set the number of threads equal to one more than the number of virtual CPUs (this will be the same as the number of physical cores if you don't have hyperthreading). Calculate the speedup and compare it to using an equal number of threads as virtual CPUs. Which performs better? Write a few sentences explaining any similarities or differences.

Submitting

Keep the same folder structure you were given. Archive in `tar.gz` format (or `zip` if you must) the `assignment-01` folder. It should contain the following in that folder: `Makefile`, `src/montecarlo.c` and `report.pdf` (you can modify the provided `report/report.tex` and create it with `make report`). After running `make` in the `assignment-01` folder it should produce two files, `bin/montecarlo` and `bin/montecarlo_parallel` in that folder. When completed, submit your archive on the course website, <http://patricklam.ca/p4p/>, after logging in.