

# **Lecture 01—Introduction**

## **ECE 459: Programming for Performance**

Patrick Lam

University of Waterloo

January 8, 2013

[Thanks to Jon Eyolfson for slides!]

# Course Website

<http://patricklam.ca/p4p/>

# Staff

## Instructor

Patrick Lam   p.lam@ece.uwaterloo.ca   DC 2597D/DC2534

## Teaching Assistants

Xavier Noumbissi   xnoumbis@uwaterloo.ca   DC 2553  
Morteza Nabavi   mnabavi@uwaterloo.ca   E5-4131

# Schedule

**Lectures:** January 8—April 4, TTh, 8:30 AM, RCH 103


**Tutorials:** January 11—April 5, F, 4:30 PM, RCH 103

**Midterm:** February 27, W, 7:00 PM—8:20PM, RCH 301/309

# Office Hours

- Suggestions?

# Recommended Textbook



Multicore Application Programming For Windows, Linux, and Oracle Solaris. Darryl Gove. Addison-Wesley, 2010.

# Goal

Make programs run faster!

# Making Programs Faster

Two main ways:

- Increase bandwidth (tasks per unit time); or
- Decrease latency (time per task).

## **Examples of bandwidth/latency:**

Network (connection speed/ping), traffic (lanes/speed)



# Our Focus

Primarily on increasing bandwidth (more tasks/unit time).

- Do tasks in parallel

Decreasing time/task usually harder, with fewer gains.

CPUs have been going towards more cores rather than raw speed.

# A Bit on Improving Latency

We won't return to these topics, but we'll touch on them now.

- Profile the code;
- Do less work;
- Be smarter; or
- Improve the hardware.

# Increasing Bandwidth: Parallelism

Some tasks are easy to run in parallel.

**Examples:** web server requests, computer graphics, brute-force searches, genetic algorithms

Others are more difficult.

**Example:** linked list traversal (why?)

# Hardware

- Use pipelining (all modern CPU do this):
  - ▶ Implement this in software by splitting a task into subtasks and running the subtasks in parallel
- Increase the number of cores/CPU's.
- Use multiple connected machines.
- Use specialized hardware, such as a GPU which contains hundreds of simple cores.

# Barriers to parallelization

- Independent tasks (“embarrassingly parallel problems”) are trivial to parallelize, but dependencies cause problems.
- Unable to start task until previous task finishes.
- May require synchronization and combination of results.
- More difficult to reason about, since execution may happen in any order.

# Limitations

- Sequential tasks in the problem will always dominate maximum performance
- Some sequential problems may be parallelizable by reformulating the implementation
- However, no matter how many processors you have, you won't be able to speed up the program as a whole (known as **Amdahl's Law**)

# Data Race

- Two processors accessing the same data.

- For example, consider the following code:

```
x = 1
```

```
print x
```

**You run it and see it prints 5**

- **Why?** Before the print, another thread wrote a new value for `x`. This is an example of a data race.

# Deadlock

Two processors trying to access a shared resource.

- Consider two processors trying to get two resources:

## Processor 1

Get Resource 1

Get Resource 2

Release Resource 2

Release Resource 1

## Processor 2

Get Resource 2

Get Resource 1

Release Resource 1

Release Resource 2

- Processor 1 gets Resource 1, then Processor 2 gets Resource 2, now they both wait for each other (**deadlock**).



# Objectives

- Implementation of parallel programming involving synchronization
- Understanding of parallel computing frameworks
- Ability to investigate software and improve its performance
- Specialized GPU programming/programming languages

# Assignments

- 1 Manual parallelization using Pthreads
- 2 Automatic parallelization and OpenMP
- 3 Application profiling and improvement (groups of 2)
- 4 GPU programming

# Breakdown

- 40% Assignments (10% each)
- 10% Midterm
- 50% Final

# Grace Days

- 4 grace days to use over the semester for late assignments
- **No mark penalty** for using grace days
- Try not to use them just because they're there

# Suggestions?

- Just let me know