Lecture 01—Introduction
ECE 459: Programming for Performance

Patrick Lam

University of Waterloo

January 8, 2013

[Thanks to Jon Eyolfson for slides!]

Course Website

http://patricklam.ca/p4p/

Staff

Instructor
Patrick Lam   p.lam@ece.uwaterloo.ca   DC 2597D/DC2534

Teaching Assistants
[censored]   censored@uwaterloo.ca   DC 2553
Morteza Nabavi   mnabavi@uwaterloo.ca   E5-4131

Schedule

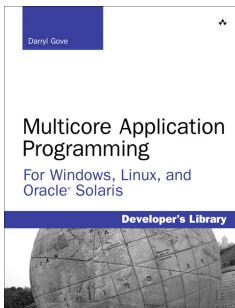**Lectures:** January 8—April 4, TTh, 8:30 AM, RCH 103

Tutorials: January 11—April 5, F, 4:30 PM, RCH 103

Midterm: February 27, W, 7:00 PM—8:20PM, RCH 301/309

Office Hours

- Suggestions?

Recommended Textbook

Multicore Application Programming For Windows, Linux, and Oracle Solaris. Darryl Gove. Addison-Wesley, 2010.

Goal

Make programs run faster!

Making Programs Faster

Two main ways:

- Increase bandwidth (tasks per unit time); or
- Decrease latency (time per task).

Examples of bandwidth/latency:
Network (connection speed/ping), traffic (lanes/speed)

Our Focus

Primarily on increasing bandwidth (more tasks/unit time).

- Do tasks in parallel

Decreasing time/task usually harder, with fewer gains.

CPUs have been going towards more cores rather than raw speed.

A Bit on Improving Latency

We won't return to these topics, but we'll touch on them now.

- Profile the code;
- Do less work;
- Be smarter; or
- Improve the hardware.

Increasing Bandwidth: Parallelism

Some tasks are easy to run in parallel.

Examples: web server requests, computer graphics, brute-force searches, genetic algorithms

Others are more difficult.

Example: linked list traversal (why?)

Hardware

- Use pipelining (all modern CPU do this):
  - Implement this in software by spliting a task into subtasks and running the subtasks in parallel

- Increase the number of cores/CPUs.

- Use multiple connected machines.

- Use specialized hardware, such as a GPU which contains hundreds of simple cores.

Barriers to parallelization

- Independent tasks ("embarrassingly parallel problems") are trivial to parallelize, but dependencies cause problems.

- Unable to start task until previous task finishes.

- May require synchronization and combination of results.

- More difficult to reason about, since execution may happen in any order.

Limitations

- Sequential tasks in the problem will always dominate maximum performance

- Some sequential problems may be parallelizable by reformulating the implementation

- However, no matter how many processors you have, you won't be able to speed up the program as a whole (known as **Amdahl's Law**)

Data Race

- Two processors accessing the same data.

- For example, consider the following code:
  ```
  x = 1
  print x
  ```
  You run it and see it prints 5

- **Why?** Before the print, another thread wrote a new value for x. This is an example of a data race.

Deadlock

Two processors trying to access a shared resource.

- Consider two processors trying to get two resources:

Processor 1	**Processor 2**
Get Resource 1	Get Resource 2
Get Resource 2	Get Resource 1
Release Resource 2	Release Resource 1
Release Resource 1	Release Resource 2

- Processor 1 gets Resource 1, then Processor 2 gets Resource 2, now they both wait for each other **(deadlock)**.

Objectives

- Implement parallel programming involving synchronization

- Describe and use parallel computing frameworks

- Ability to investigate software and improve its performance

- Use and understand specialized GPU programming/programming languages

Assignments

1 Manual parallelization using Pthreads

2 Automatic parallelization and OpenMP

3 Application profiling and improvement (groups of 2)

4 GPU programming

Breakdown

- 40% Assignments (10% each)

- 10% Midterm

- 50% Final

Grace Days

- 4 grace days to use over the semester for late assignments

- **No mark penalty** for using grace days

- Try not to use them just because they're there

Suggestions?

- Just let me know