

Lecture 04—Pthreads and Simple Locks (v2)

ECE 459: Programming for Performance

January 17, 2013

Background

Recall the difference between `processes` and `threads`:

- Threads are basically light-weight processes which piggy-back on processes' address space.

Traditionally (pre-Linux 2.6) you had to use `fork` (for processes) and `clone` (for threads)

History

`clone` is not POSIX compliant.

Developers mostly used `fork` in the past, which creates a new process.

- Drawbacks?
- Benefits?

Benefit: `fork` is Safer and More Secure Than Threads

- ❶ Each process has its own virtual address space:
 - ▶ Memory pages are not copied, they are copy-on-write—
 - ▶ Therefore, uses less memory than you would expect.
- ❷ Buffer overruns or other security holes do not expose other processes.
- ❸ If a process crashes, the others can continue.

Example: In the Chrome browser, each tab is a separate process.

Drawback of Processes: Threads are Easier and Faster

- Interprocess communication (IPC) is more complicated and slower than interthread communication.
 - ▶ Need to use operating system utilities (pipes, semaphores, shared memory, etc) instead of thread library (or just memory reads and writes).
- Processes have much higher startup, shutdown and synchronization cost.
- And, **Pthreads** fix the issues of `clone` and provide a uniform interface for most systems (**focus of Assignment 1**).

Appropriate Time to Use Processes

If your application is like this:

- Mostly independent tasks, with little or no communication.
- Task startup and shutdown costs are negligible compared to overall runtime.
- Want to be safer against bugs and security holes.

Then processes are the way to go.

For performance reasons, along with ease and consistency, we'll use [Pthreads](#).

fork Usage Example (OS refresher)

```
pid = fork();  
if (pid < 0) {  
    fork_error_function();  
} else if (pid == 0) {  
    child_function();  
} else {  
    parent_function();  
}
```

`fork` produces a second copy of the calling process, which starts execution after the call.

The only difference between the copies is the return value: the parent gets the pid of the child, while the child gets 0.

Threads Offer a Speedup of 6.5 over fork

Here's a benchmark between `fork` and `Pthreads` on a laptop, creating and destroying 50,000 threads:

```
jon@riker examples master % time ./create_fork
0.18s user 4.14s system 34% cpu 12.484 total
jon@riker examples master % time ./create_pthread
0.73s user 1.29s system 107% cpu 1.887 total
```

Clearly `Pthreads` incur much lower overhead than `fork`.

Assumptions

First, we'll see how to use threads on “embarrassingly parallel problems”.

- mostly-independent sub-problems (little synchronization); and
- strong locality (little communication).

Later, we'll see:

- which problems can be parallelized ([dependencies](#))
- alternative parallelization patterns
(right now, just use one thread per sub-problem)

POSIX Threads

- Available on most systems
- Windows has Pthreads Win32, but I wouldn't use it; use Linux for this course
- API available by `#include <pthread.h>`
- Compile with pthread flag
(`gcc -pthread prog.c -o prog`)

Creating Threads

```
int pthread_create(pthread_t* thread ,  
                  const pthread_attr_t* attr ,  
                  void* (*start_routine)(void*) ,  
                  void* arg );
```

thread: creates a handle to a thread at pointer location

attr: thread attributes (NULL for defaults, more details later)

start_routine: function to start execution

arg: value to pass to start_routine

returns 0 on success, error number otherwise
(contents of *thread are undefined)

Creating Threads—Example

```
#include <pthread.h>
#include <stdio.h>

void* run(void*) {
    printf("In run\n");
}

int main() {
    pthread_t thread;
    pthread_create(&thread, NULL, &run, NULL);
    printf("In main\n");
}
```

Simply creates a thread and terminates
(usage isn't really right, as we'll see.)

Waiting for Threads

```
int pthread_join(pthread_t thread,  
                 void** retval)
```

thread: wait for this thread to terminate (thread must be joinable).

retval: stores exit status of thread (set by `pthread_exit`) to the location pointed by `*retval`. If cancelled, returns `PTHREAD_CANCELED`. `NULL` is ignored.

returns 0 on success, error number otherwise.

Only call this one time per thread! Multiple calls on the same thread leads to undefined behaviour.

Waiting for Threads—Example

```
#include <pthread.h>
#include <stdio.h>

void* run(void*) {
    printf("In run\n");
}

int main() {
    pthread_t thread;
    pthread_create(&thread, NULL, &run, NULL);
    printf("In main\n");
    pthread_join(thread, NULL);
}
```

This now waits for the newly created thread to terminate.

Passing Data to Threads... Wrongly

Consider this snippet:

```
int i;  
for (i = 0; i < 10; ++i)  
    pthread_create(&thread[i], NULL, &run, (void*)&i);
```

This is a **terrible** idea. Why?

Passing Data to Threads... Wrongly

Consider this snippet:

```
int i;  
for (i = 0; i < 10; ++i)  
    pthread_create(&thread[i], NULL, &run, (void*)&i);
```

This is a **terrible** idea. Why?

- ❶ The value of `i` will probably change before the thread executes
- ❷ The memory for `i` may be out of scope, and therefore invalid by the time the thread executes

Passing Data to Threads

What about:

```
int i;  
for (i = 0; i < 10; ++i)  
    pthread_create(&thread[i], NULL, &run, (void*)i);  
  
...  
  
void* run(void* arg) {  
    int id = (int)arg;
```

This is suggested in the book, but should carry a warning:

Passing Data to Threads

What about:

```
int i;  
for (i = 0; i < 10; ++i)  
    pthread_create(&thread[i], NULL, &run, (void*)i);  
  
...  
  
void* run(void* arg) {  
    int id = (int)arg;
```

This is suggested in the book, but should carry a warning:

- Beware size mismatches between arguments: no guarantee that a pointer is the same size as an int, so your data may overflow.
- Sizes of data types change between systems. For maximum portability, just use pointers you got from `malloc`.

Detached Threads

Joinable threads (the default) wait for someone to call `pthread_join` before they release their resources.

Detached threads release their resources when they terminate, without being joined.

```
int pthread_detach(pthread_t thread);
```

thread: marks the thread as detached

returns 0 on success, error number otherwise.

Calling `pthread_detach` on an already detached thread results in undefined behaviour.

Thread Termination

```
void pthread_exit(void *retval);
```

retval: return value passed to function that calls
`pthread_join`

`start_routine` returning is equivalent to calling `pthread_exit`
with that return value;

`pthread_exit` is called implicitly when the `start_routine` of
a thread returns.

Attributes

By default, threads are *joinable* on Linux, but a more portable way to know what you're getting is to set thread attributes.

You can change:

- Detached or joinable state
- Scheduling inheritance
- Scheduling policy
- Scheduling parameters
- Scheduling contention scope
- Stack size
- Stack address
- Stack guard (overflow) size

Attributes—Example

```
size_t stacksize;  
pthread_attr_t attributes;  
pthread_attr_init(&attributes);  
pthread_attr_getstacksize(&attributes, &stacksize);  
printf("Stack size = %i\n", stacksize);  
pthread_attr_destroy(&attributes);
```

Running this on a laptop produces:

```
jon@riker examples master % ./stack_size  
Stack size = 8388608
```

Setting a thread state to joinable:

```
pthread_attr_setdetachstate(&attributes,  
                           PTHREAD_CREATE_JOINABLE);
```

Detached Threads: Warning!

```
#include <pthread.h>
#include <stdio.h>

void* run(void*) {
    printf("In run\n");
}

int main() {
    pthread_t thread;
    pthread_create(&thread, NULL, &run, NULL);
    pthread_detach(thread);
    printf("In main\n");
}
```

When I run it, it just prints "In main", why?

Detached Threads: Solution to Problem

```
#include <pthread.h>
#include <stdio.h>

void* run(void*) {
    printf("In run\n");
}

int main() {
    pthread_t thread;
    pthread_create(&thread, NULL, &run, NULL);
    pthread_detach(thread);
    printf("In main\n");
    pthread_exit(NULL); // This waits for all detached
                        // threads to terminate
}
```

Make the final call `pthread_exit` if you have any detached threads.

Three Useful Routines

```
pthread_t  pthread_self(void);

int  pthread_equal(pthread_t t1, pthread_t t2);

int  pthread_once(pthread_once_t* once_control ,
                  void (*init_routine)(void));
pthread_once_t  once_control = PTHREAD_ONCE_INIT;
```

`pthread_self` returns the handle of the currently running thread.

Use `pthread_equal` if you're comparing 2 threads.

If you want to run a section of code once, you need `pthread_once` (it's well-named). It will run only once per `once_control`.

Threading Challenges

- Be aware of scheduling (you can also set affinity with pthreads on Linux).
- Make sure the libraries you use are **thread-safe**:
 - ▶ Means that the library protects its shared data.
- glibc reentrant functions are also safe: a program can have more than one thread calling these functions concurrently.
- **Example:** In Assignment 1, we'll use `rand_r`, not `rand`.

Mutual Exclusion

Mutexes are the most basic type of synchronization.

- Only one thread can access code protected by a mutex at a time.
- All other threads must wait until the mutex is free before they can execute the protected code.

Creating Mutexes—Example

```
pthread_mutex_t m1 = PTHREAD_MUTEX_INITIALIZER;  
pthread_mutex_t m2;  
  
pthread_mutex_init(&m2, NULL);  
...  
pthread_mutex_destroy(&m1);  
pthread_mutex_destroy(&m2);
```

- Two ways to initialize mutexes: statically and dynamically
- If you want to include attributes, you need to use the dynamic version

Mutex Attributes

- **Protocol**: specifies the protocol used to prevent priority inversions for a mutex
- **Prioceiling**: specifies the priority ceiling of a mutex
- **Process-shared**: specifies the process sharing of a mutex

You can specify a mutex as *process shared* so that you can access it between processes. In that case, you need to use shared memory and `mmap`, which we won't get into.

Using Mutexes: Example

```
// code
pthread_mutex_lock(&m1);
// protected code
pthread_mutex_unlock(&m1);
// more code
```

- Everything within the lock and unlock is protected.
- Be careful to avoid deadlocks if you are using multiple mutexes.
- Also you can use `pthread_mutex_trylock`, if needed.

Data Race Example

Recall that **dataraces** occur when two concurrent actions access the same variable and at least one of them is a **write**

```
...
static int counter = 0;

void* run(void* arg) {
    for (int i = 0; i < 100; ++i) {
        ++counter;
    }
}

int main(int argc, char *argv[])
{
    // Create 8 threads
    // Join 8 threads
    printf("counter = %i\n", counter);
}
```

Is there a datarace in this example? If so, how would we fix it?

Example Problem Solution

```
...
static pthread_mutex_t mutex = PTHREAD_MUTEX_INITIALIZER;
static int counter = 0;

void* run(void* arg) {
    for (int i = 0; i < 100; ++i) {
        pthread_mutex_lock(&mutex);
        ++counter;
        pthread_mutex_unlock(&mutex);
    }
}

int main(int argc, char *argv[])
{
    // Create 8 threads
    // Join 8 threads
    pthread_mutex_destroy(&mutex);
    printf("counter = %i\n", counter);
}
```