

Lecture 17—Midterm Results, A3 Discussion, Profiling

ECE 459: Programming for Performance

March 12, 2013

Previous Lecture


Midterm Solutions

Midterm Results

Note: I noticed that Morteza sometimes forgot to give points for Q1 (iii). Double-check your exam.

Raw mean: 69.96%

Raw median: 69.4%



Scaling

We can see two clumps:

“A”: 75+

“B”: 60–70

Here's the formula I used:

$$P_s = 5 + \frac{100}{90} M_r.$$

P_s is your midterm mark, in percent.

M_r is the raw mark.

Part I

About Assignment 3

Assignment Problem: Beier-Neely image morphing



“an image processing technique typically used as an animation tool for the metamorphosis of one image to another.”

Thaddeus Beier and Shawn Neely. “Feature-Based Image Metamorphosis”. SIGGRAPH 1992, pp. 35–42.

Assignment Problem: Beier-Neely image morphing



“an image processing technique typically used as an animation tool for the metamorphosis of one image to another.”

Thaddeus Beier and Shawn Neely. “Feature-Based Image Metamorphosis”. SIGGRAPH 1992, pp. 35–42.

Assignment Problem: Beier-Neely image morphing



“an image processing technique typically used as an animation tool for the metamorphosis of one image to another.”

Thaddeus Beier and Shawn Neely. “Feature-Based Image Metamorphosis”. SIGGRAPH 1992, pp. 35–42.

Introduction

Image morphing:

computes intermediate images between source and dest.

More complicated than cross-dissolving between pictures
(interpolating pixel colours).

Morphing = warping + cross-dissolving.

Domain Discussion Omitted

Sorry, I had no time to learn and explain how code works.

We'll see it again on Thursday.

Algorithms

All of the C++ built-in algorithms work with anything using the standard container interface.

- `max_element`—returns an iterator with the largest element; you can use your own comparator function.
- `min_element`—same as above, but the smallest element.
- `sort`—sorts a container; you can use your own comparator function.
- `upper_bound`—returns an iterator to the first element greater than the value; only works on a **sorted** container (if the default comparator isn't used, you have to use the same one used to sort the container).
- `random_shuffle`—does `n` random swaps of the elements in the container

Some Hurdles

If you've worked with C++ before, you probably know the awful compiler messages and pages of template expansions.

- You can use `clang` if you have a compiler error, and let it show you the error instead `-std=c++11`
- For the profiler messages, it might get pretty bad. Look for one of the main functions, or if it's a weird name, look where it's called from.
- You can use Google Perf Tools to help break it down—more fine grained.

Example Profiler Function Output

```
[32] std::_Hashtable<unsigned int, std::pair<unsigned
      int const, std::unordered_map<unsigned int,
      double, std::hash<unsigned int>, std::equal_to<
      unsigned int>, std::allocator<std::pair<unsigned
      int const, double>>>>, std::allocator<std::
      pair<unsigned int const, std::unordered_map<
      unsigned int, double, std::hash<unsigned int>,
      std::equal_to<unsigned int>, std::allocator<std::
      pair<unsigned int const, double>>>>>>, std::
      _Select1st<std::pair<unsigned int const, std::
      unordered_map<unsigned int, double, std::hash<
      unsigned int>, std::equal_to<unsigned int>, std::
      allocator<std::pair<unsigned int const, double>>>
      >>>, std::equal_to<unsigned int>, std::hash<
      unsigned int>, std::__detail::_Mod_range_hashing,
      std::__detail::_Default_ranged_hash, std::
      __detail::_Prime_rehash_policy, false, false,
      true>::clear()
```

is actually `distance_map.clear()` (from last year's assignment), which is automatically called by the destructor.

Things You Can Do

Well, it's a basic implementation, so there should be a lot you can do.

- You can introduce threads, using pthreads (or C++11 threads, if you're feeling lucky), OpenMP, or whatever you want.
- Play around with compiler options.
- Use better algorithms or data structures (maybe Qt is inefficient!)
- The list goes on and on.

Things You Need to Do

Profile!

- Keep your inputs constant between all profiling results so they're comparable.
- Baseline profile with no changes.
- You will pick your two best performance changes to add to the report.
 - ▶ You will include a profiling report before the change and just after the change (and only that change!)
 - ▶ More specific instructions in the handout.
- There may or may not be overlap between the baseline and the baseline for each change.
- My recommendation: use your initial baseline as the “before” for your first change, and the “after” of the first change for the baseline of your second change.
- Whatever you choose, it should be convincing.

Things To Notice

Your program will be run on ece459-1 (or equivalent), with a 10 second limit.

We will have some type of leaderboard, so the earlier you have some type of submission, the better.

A Word

This assignment should be **enjoyable**.

Part II

Profiling

Introduction to Profiling

So far we've been looking at small problems.

Must **profile** to see what takes time in a large program.

Two main outputs:

- flat;
- call-graph.

Two main data gathering methods:

- statistical;
- instrumentation.

Profiler Outputs

Flat Profiler:

- Only computes the average time in a particular function.
- Does not include other (useful) information, like callees.

Call-graph Profiler:

- Computes call times.
- Reports frequency of function calls.
- Gives a call graph: who called what function?

Data Gathering Methods

Statistical:

Mostly, take samples of the system state, that is:

- every 2ns, check the system state.
- will cause some slowdown, but not much.

Instrumentation:

Add additional instructions at specified program points:

- can do this at compile time or run time (expensive);
- can instrument either manually or automatically;
- like conditional breakpoints.

Guide to Profiling

When writing large software projects:

- First, write clear and concise code.
Don't do any premature optimizations—focus on correctness.
- Profile to get a baseline of your performance:
 - ▶ allows you to easily track any performance changes;
 - ▶ allows you to re-design your program before it's too late.

Focus your optimization efforts on the code that matters.

Things to Look For

Good signs:

- Time is spent in the right part of the system.
- Most time should not be spent handling errors; in non-critical code; or in exceptional cases.
- Time is not unnecessarily spent in the operating system.

gprof introduction

Statistical profiler, plus some instrumentation for calls.

Runs completely in user-space.

Only requires a compiler.

gprof usage

Use the `-pg` flag with `gcc` when compiling and linking.

Run your program as you normally would.

- Your program will now create a `gmon.out` file.

Use `gprof` to interpret the results: `gprof <executable>`.

gprof example

A program with 100 million calls to two math functions.

```
int main() {  
    int i, x1=10, y1=3, r1=0;  
    float x2=10, y2=3, r2=0;  
  
    for( i=0; i < 100000000; i++) {  
        r1 += int_math(x1, y1);  
        r2 += float_math(y2, y2);  
    }  
}
```

- Looking at the code, we have no idea what takes longer.
- Probably would guess floating point math taking longer.
- (Overall, silly example.)

Example (Integer Math)

```
int int_math(int x, int y){
    int r1;
    r1=int_power(x,y);
    r1=int_math_helper(x,y);
    return r1;
}

int int_math_helper(int x, int y){
    int r1;
    r1=x/y*int_power(y,x)/int_power(x,y);
    return r1;
}

int int_power(int x, int y){
    int i, r;
    r=x;
    for(i=1;i<y;i++){
        r=r*x;
    }
    return r;
}
```

Example (Float Math)

```
float float_math(float x, float y) {  
    float r1;  
    r1=float_power(x,y);  
    r1=float_math_helper(x,y);  
    return r1;  
}  
  
float float_math_helper(float x, float y) {  
    float r1;  
    r1=x/y*float_power(y,x)/float_power(x,y);  
    return r1;  
}  
  
float float_power(float x, float y){  
    float i, r;  
    r=x;  
    for(i=1;i<y;i++) {  
        r=r*x;  
    }  
    return r;  
}
```

Flat Profile

When we run the program and look at the profile, we see:

Flat profile :

Each sample counts as 0.01 seconds.

% time	cumulative seconds	self seconds	calls	self ns/call	total ns/call	name
32.58	4.69	4.69	300000000	15.64	15.64	int_power
30.55	9.09	4.40	300000000	14.66	14.66	float_power
16.95	11.53	2.44	100000000	24.41	55.68	int_math_helper
11.43	13.18	1.65	100000000	16.46	45.78	float_math_helper
4.05	13.76	0.58	100000000	5.84	77.16	int_math
3.01	14.19	0.43	100000000	4.33	64.78	float_math
2.10	14.50	0.30				main

- One function per line.
- **% time:** the percent of the total execution time in this function.
- **self:** seconds in this function.
- **cumulative:** sum of this function's time + any above it in table.

Flat Profile

Flat profile:

Each sample counts as 0.01 seconds.

% time	cumulative seconds	self seconds	calls	self ns/call	total ns/call	name
32.58	4.69	4.69	300000000	15.64	15.64	int_power
30.55	9.09	4.40	300000000	14.66	14.66	float_power
16.95	11.53	2.44	100000000	24.41	55.68	int_math_helper
11.43	13.18	1.65	100000000	16.46	45.78	float_math_helper
4.05	13.76	0.58	100000000	5.84	77.16	int_math
3.01	14.19	0.43	100000000	4.33	64.78	float_math
2.10	14.50	0.30				main

- **calls:** number of times this function was called
- **self ns/call:** just self nanoseconds / calls
- **total ns/call:** average time for function execution, including any other calls the function makes

Call Graph Example (1)

After the flat profile gives you a feel for which functions are costly, you can get a better story from the call graph.

index	% time	self	children	called	name
<spontaneous>					
[1]	100.0	0.30	14.19		main [1]
		0.58	7.13	100000000/100000000	int_math [2]
		0.43	6.04	100000000/100000000	float_math [3]
<hr/>					
[2]	53.2	0.58	7.13	100000000/100000000	main [1]
		0.58	7.13	100000000	int_math [2]
		2.44	3.13	100000000/100000000	int_math_helper [4]
		1.56	0.00	100000000/300000000	int_power [5]
<hr/>					
[3]	44.7	0.43	6.04	100000000/100000000	main [1]
		0.43	6.04	100000000	float_math [3]
		1.65	2.93	100000000/100000000	float_math_helper [6]
		1.47	0.00	100000000/300000000	float_power [7]

Reading the Call Graph

The line with the index is the current function being looked at
(primary line).

- Lines above are functions which called this function.
- Lines below are functions which were called by this function (children).

Primary Line

- **time:** total percentage of time spent in this function and its children
- **self:** same as in flat profile
- **children:** time spent in all calls made by the function
 - ▶ should be equal to self + children of all functions below

Reading Callers from Call Graph

Callers (functions above the primary line)

- **self:** time spent in primary function, when called from current function.
- **children:** time spent in primary function's children, when called from current function.
- **called:** number of times primary function was called from current function / number of nonrecursive calls to primary function.

Reading Callees from Call Graph

Callees (functions below the primary line)

- **self:** time spent in current function when called from primary.
- **children:** time spent in current function's children calls when called from primary.
 - ▶ $\text{self} + \text{children}$ is an estimate of time spent in current function when called from primary function.
- **called:** number of times current function was called from primary function / number of nonrecursive calls to current function.

Call Graph Example (2)

index	% time	self	children	called	name
[4]	38.4	2.44	3.13	100000000/100000000	int_math [2]
		2.44	3.13	100000000	int_math_helper [4]
		3.13	0.00	200000000/300000000	int_power [5]
[5]	32.4	1.56	0.00	100000000/300000000	int_math [2]
		3.13	0.00	200000000/300000000	int_math_helper [4]
		4.69	0.00	300000000	int_power [5]
[6]	31.6	1.65	2.93	100000000/100000000	float_math [3]
		1.65	2.93	100000000	float_math_helper [6]
		2.93	0.00	200000000/300000000	float_power [7]
[7]	30.3	1.47	0.00	100000000/300000000	float_math [3]
		2.93	0.00	200000000/300000000	float_math_helper [6]
		4.40	0.00	300000000	float_power [7]

We can now see where most of the time comes from, and pinpoint any locations that make unexpected calls, etc.

This example isn't too exciting; we could simplify the math.

Part III

gperftools

Introduction to gperftools

Google Performance Tools include:

- CPU profiler.
- Heap profiler.
- Heap checker.
- Faster `malloc`.

We'll mostly use the CPU profiler:

- purely statistical sampling;
- no recompilation; at most linking; and
- built-in visual output.

Google Perf Tools profiler usage

You can use the profiler without any recompilation.

- Not recommended—worse data.

```
LD_PRELOAD="/usr/lib/libprofiler.so" \  
CPUPROFILE=test.prof ./test
```

The other option is to link to the profiler:

- `-lprofiler`

Both options read the `CPUPROFILE` environment variable:

- states the location to write the profile data.

Other Usage

You can use the profiling library directly as well:

- `#include <gperftools/profiler.h>`

Bracket code you want profiled with:

- `ProfilerStart()`
- `ProfilerEnd()`

You can change the sampling frequency with the `CPUPROFILE_FREQUENCY` environment variable.

- **Default value:** 100

pprof Usage

Like gprof, it will analyze profiling results.

```
% pprof test test.prof
    Enters "interactive" mode
% pprof --text test test.prof
    Outputs one line per procedure
% pprof --gv test test.prof
    Displays annotated call-graph via 'gv'
% pprof --gv --focus=Mutex test test.prof
    Restricts to code paths including a .*Mutex.* entry
% pprof --gv --focus=Mutex --ignore=string test test.prof
    Code paths including Mutex but not string
% pprof --list=getdir test test.prof
    (Per-line) annotated source listing for getdir()
% pprof --disasm=getdir test test.prof
    (Per-PC) annotated disassembly for getdir()
```

Can also output dot, ps, pdf or gif instead of gv.

Text Output

Similar to the flat profile in gprof

```
jon@riker examples master % pprof --text test test.prof
Using local file test.
Using local file test.prof.
Removing killpg from all stack traces.
Total: 300 samples
```

95	31.7%	31.7%	102	34.0%	int_power
58	19.3%	51.0%	58	19.3%	float_power
51	17.0%	68.0%	96	32.0%	float_math_helper
50	16.7%	84.7%	137	45.7%	int_math_helper
18	6.0%	90.7%	131	43.7%	float_math
14	4.7%	95.3%	159	53.0%	int_math
14	4.7%	100.0%	300	100.0%	main
0	0.0%	100.0%	300	100.0%	__libc_start_main
0	0.0%	100.0%	300	100.0%	_start

Text Output Explained

Columns, from left to right:

Number of checks (samples) in this function.

Percentage of checks in this function.

- Same as **time** in gprof.

Percentage of checks in the functions printed so far.

- Equivalent to **cumulative** (but in %).


Number of checks in this function and its callees.

Percentage of checks in this function and its callees.

Function name.

Graphical Output

a.out
Total samples: 300
Focusing on: 300
Dropped nodes with ≤ 1 abs(samples)
Dropped edges with ≤ 0 samples



Graphical Output Explained

Output was too small to read on the slide.

- Shows the same numbers as the text output.
- Directed edges denote function calls.
- Note: number of samples in callees =
number in “this function + callees” -
number in “this function”.
- **Example:**
float_math_helper, 51 (local) of 96 (cumulative)
 $96 - 51 = 45$ (callees)
 - ▶ callee int_power = 7 (bogus)
 - ▶ callee float_power = 38
 - ▶ callees total = 45

Things You May Notice

Call graph is not exact.

- In fact, it shows many bogus relations which clearly don't exist.
- For instance, we know that there are no cross-calls between `int` and `float` functions.

As with `gprof`, optimizations will change the graph.

You'll probably want to look at the text profile first, then use the `-focus` flag to look at individual functions.

Summary

- Talked about midterm and A3 (more on Thursday).
- Saw how to use gprof
(one option for Assignment 3).
- Profile early and often.
- Make sure your profiling shows what you expect.
- We'll see other profilers we can use as well:
 - ▶ OProfile
 - ▶ Valgrind
 - ▶ AMD CodeAnalyst
 - ▶ Perf