

Lecture 18—More A3, More Profiling Tools

ECE 459: Programming for Performance

March 14, 2013

Part I

More A3 Discussion

Morphing = Warping + Cross-Dissolving



⇓ warp



dissolve
⇒



dissolve
⇐



warp ⇓




Warping Examples: Single-Line Warp

Next few figures are from the original paper.

Another reference:

http://www.cs.unc.edu/~lazebnik/research/fall08/qi_mo.pdf.



$$u = \frac{(X - P) \cdot (Q - P)}{\|Q - P\|^2} \quad (1)$$


$$v = \frac{(X - P) \cdot \text{Perpendicular}(Q - P)}{\|Q - P\|} \quad (2)$$

$$\textcircled{X'} = P' + u \cdot (Q' - P') + \frac{v \cdot \text{Perpendicular}(Q' - P')}{\|Q' - P'\|} \quad (3)$$

For each point in the destination, use the corresponding transformed point in the source.

Warping Examples: Multiple-Line Warp



$$D_i = X_i' - X_i$$



$$weight = \left(\frac{length^p}{(a + dist)} \right)^h$$

Generalize the definition of “corresponding point” for multiple lines.

Warped Image



Parameters: a, b, p

The warping algorithm takes three parameters:

- a : smoothness of warping
(a near 0 means that lines go to lines)
- b : how relative strength of lines falls off with distance.
large means every pixel only affected by nearest line;
0 means each pixel affected by all lines equally.
suggested range: $[0.5, 2]$
- p : relationship between length of line and strength.
0 means all lines have same weight;
1 means longer lines have greater relative weight;
suggested range: $[0, 1]$.

Algorithm Pseudocode

```
for each pixel  $X$  in the destination:  
  DSUM  $\leftarrow (0, 0)$   
  weightsum  $\leftarrow 0$   
  for each line  $P_i Q_i$ :  
    calculate  $u, v$  based on  $P_i Q_i$   
    calculate  $X'_i$  based on  $u, v$  and  $P'_i Q'_i$   
    calculate displacement  $D_i = X'_i - X_i$  for this line  
    dist  $\leftarrow$  shortest distance from  $X$  to  $P_i Q_i$   
    weight  $\leftarrow (\text{length}^p / (a + \text{dist}))^b$   
    DSUM  $+= D_i \times \text{weight}$   
    weightsum  $+= \text{weight}$   
   $X' = X + \text{DSUM} / \text{weightsum}$   
  destinationImage( $X$ )  $\leftarrow$  sourceImage( $X'$ )
```


Introduction

- The code is more or less a direct translation of the high level functions.
- The language should not be the main hurdle, if there's anything you don't understand about the provided code, feel free to talk to me.
- Code uses standard library functions plus Qt types.

Built-in Data Structures

- QPoint
- QVector2D
- QImage

This is a fairly straightforward computation.

Functions

All of the code that you need to deal with is in `model.cpp`.

For your purposes, this file is called from `test_harness.cpp`; it is also called from `window.cpp`, the interactive front-end.

`model` contains three methods:

- `prepStraightLine`: draws the lines and computes auxiliary lines;
- `commonPrep`: draws auxiliary lines;
- `morph`: carries out the actual morphing algorithm.

I refactored `model` out of the initial `window.cpp`, and may have made some mistakes. If you find them, fix them.

Notes

I plan to add some more test cases and tweak the parameters shortly.

You'll probably want to refactor `morph()` to get useful profiling information from it.

You can remove code if you can prove it useless (using tests and text, which you should have anyway.)

Part II

System profiling: oprofile, perf, DTrace,
WAIT

Introduction: oprofile

`http://oprofile.sourceforge.net`

Sampling-based tool.

Uses CPU performance counters.

Tracks currently-running function;
records profiling data for every application run.

Can work system-wide (across processes).

Technology: Linux Kernel Performance Events
(formerly a Linux kernel module).

Setting up oprofile

Must run as root to use system-wide, otherwise can use per-process.

```
% sudo opcontrol \  
    --vmlinux=/usr/src/linux-3.2.7-1-ARCH/vmlinux  
% echo 0 | sudo tee /proc/sys/kernel/nmi_watchdog  
% sudo opcontrol --start  
Using default event: CPU_CLK_UNHALTED:100000:0:1:1  
Using 2.6+ OProfile kernel interface.  
Reading module info.  
Using log file /var/lib/oprofile/samples/oprofiled.log  
Daemon started.  
Profiler running.
```

Per-process:

```
[plam@lynch nm-morph]$ operf ./test_harness  
operf: Profiler started  
  
Profiling done.
```

oprofile Usage (1)

Pass your executable to opreport.

```
% sudo opreport -l ./test
CPU: Intel Core/i7 , speed 1595.78 MHz (estimated)
Counted CPU_CLK_UNHALTED events (Clock cycles when not
halted) with a unit mask of 0x00 (No unit mask) count 100000
samples  %          symbol name
7550      26.0749   int_math_helper
5982      20.6596   int_power
5859      20.2348   float_power
3605      12.4504   float_math
3198      11.0447   int_math
2601       8.9829   float_math_helper
160        0.5526   main
```

If you have debug symbols (-g) you could use:

```
% sudo opannotate --source \
--output-dir=/path/to/annotated-source /path/to/mybinary
```


oprofile Usage (2)

Use `opreport` by itself for a whole-system view.
You can also reset and stop the profiling.

```
% sudo opcontrol --reset  
Signalling daemon... done  
% sudo opcontrol --stop  
Stopping profiling.
```

Perf: Introduction

`https://perf.wiki.kernel.org/index.php/Tutorial`

Interface to Linux kernel built-in sampling-based profiling.

Per-process, per-CPU, or system-wide.

Can even report the cost of each line of code.

Perf: Usage Example

On the Assignment 3 code:

```
[plam@lynch nm-morph]$ perf stat ./test_harness
```

```
Performance counter stats for './test_harness':
```

| | | | | | |
|----------------|-------------------------|---|---------|-------------------------|----------|
| 6562.501429 | task-clock | # | 0.997 | CPU's utilized | |
| 666 | context-switches | # | 0.101 | K/sec | |
| 0 | cpu-migrations | # | 0.000 | K/sec | |
| 3,791 | page-faults | # | 0.578 | K/sec | |
| 24,874,267,078 | cycles | # | 3.790 | GHz | [83.32%] |
| 12,565,457,337 | stalled-cycles-frontend | # | 50.52% | frontend cycles idle | [83.31%] |
| 5,874,853,028 | stalled-cycles-backend | # | 23.62% | backend cycles idle | [66.63%] |
| 33,787,408,650 | instructions | # | 1.36 | insns per cycle | |
| | | # | 0.37 | stalled cycles per insn | [83.32%] |
| 5,271,501,213 | branches | # | 803.276 | M/sec | [83.38%] |
| 155,568,356 | branch-misses | # | 2.95% | of all branches | [83.36%] |
| 6.580225847 | seconds time elapsed | | | | |

Perf: Source-level Analysis

perf can tell you which instructions are taking time, or which lines of code.

Compile with `-ggdb` to enable source code viewing.

```
% perf record ./test_harness  
% perf annotate
```

`perf annotate` is interactive. Play around with it.

DTrace: Introduction

`http://queue.acm.org/detail.cfm?id=1117401`

Intrumentation-based tool.

System-wide.

Meant to be used on production systems. (Eh?)

DTrace: Introduction

`http://queue.acm.org/detail.cfm?id=1117401`

Instrumentation-based tool.

System-wide.

Meant to be used on production systems. (Eh?)

(Typical instrumentation can have a slowdown of 100x (Valgrind).)

Design goals:

- 1 No overhead when not in use;
- 2 Guarantee safety—must not crash
(strict limits on expressiveness of probes).

DTrace: Operation

How does DTrace achieve 0 overhead?

- only when activated, dynamically rewrites code by placing a branch to instrumentation code.

Uninstrumented: runs as if nothing changed.

Most instrumentation: at function entry or exit points.
You can also instrument kernel functions, locking,
instrument-based on other events.

Can express sampling as instrumentation-based events also.

DTrace Example

You write this:

```
syscall::read:entry {  
    self->t = timestamp;  
}  
  
syscall::read:return  
/self->t/ {  
    printf("%d/%d spent %d nsecs in read\n"  
          pid, tid, timestamp - self->t);  
}
```

`t` is a thread-local variable.

This code prints how long each call to `read` takes, along with context.

To ensure safety, DTrace limits what you write; e.g. no loops.

- (Hence, no infinite loops!)

Other Tools

AMD CodeAnalyst—based on oprofile; leverages AMD processor features.

WAIT

- IBM's tool tells you what operations your JVM is waiting on while idle.
- Non-free and not available.

Other Tools

AMD CodeAnalyst—based on oprofile.

WAIT

- IBM's tool tells you what operations your JVM is waiting on while idle.
- Non-free and not available.

WAIT: Introduction

Built for production environments.

Specialized for profiling JVMs, uses JVM hooks to analyze idle time.

Sampling-based analysis; infrequent samples (1–2 per minute!)

WAIT: Operation

At each sample: records each thread's state,

- call stack;
- participation in system locks.

Enables WAIT to compute a “wait state” (using expert-written rules):

what the process is currently doing or waiting on, e.g.

- disk;
- GC;
- network;
- blocked;
- etc.

WAIT: Workflow

You:

- run your application;
- collect data (using a script or manually); and
- upload the data to the server.

Server provides a report.

- You fix the performance problems.

Report indicates processor utilization (idle, your application, GC, etc); runnable threads; waiting threads (and why they are waiting); thread states; and a stack viewer.

Paper presents 6 case studies where WAIT identified performance problems: deadlocks, server underloads, memory leaks, database bottlenecks, and excess filesystem activity.

Other Profiling Tools

Profiling: Not limited to C/C++, or even code.

You can profile Python using `cProfile`; standard profiling technology.

Google's Page Speed Tool: profiling for web pages—how can you make your page faster?

- reducing number of DNS lookups;
- leveraging browser caching;
- combining images;
- plus, traditional JavaScript profiling.

Summary

- More Assignment 3 discussion
- System profiling tools:
 - oprofile, DTrace, WAIT, perf

Future Lectures

- OpenMPI
- OpenCL
- Hadoop MapReduce
- Software Transactional Memory
- Early Phase Termination, Big Data

If there's anything else you want to see in this course, let me know.